



# Flash-Friendly File System (F2FS)

Presented at KLF 2012

October 12, 2012

Jaegeuk Kim

S/W Development Team

Memory Division, Samsung Electronics Co., Ltd.



TURN ON TOMORROW


# Agenda



- Introduction
- Log-structured File System
- Design Issues
- Design of F2FS
- Performance Evaluation
- Summary


# Introduction

- NAND Flash-based Storage Devices
  - SSD for PC and server systems
  - eMMC for mobile systems
  - SD card for consumer electronics
- The Rise of SSDs
  - Much faster than HDDs
  - Low power consumption



Source: March 30th, 2012 by Avram Piltch, LAPTOP Online Editorial Director

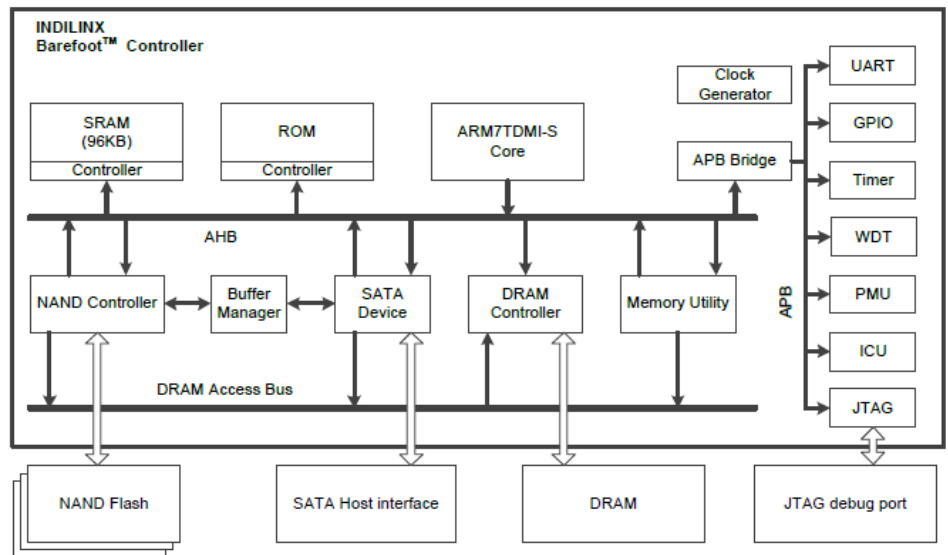
Figure-3 2008-2013 Solid-State Drive Market Forecast



Source: DRAmEXchange, Jan., 2012

# Introduction

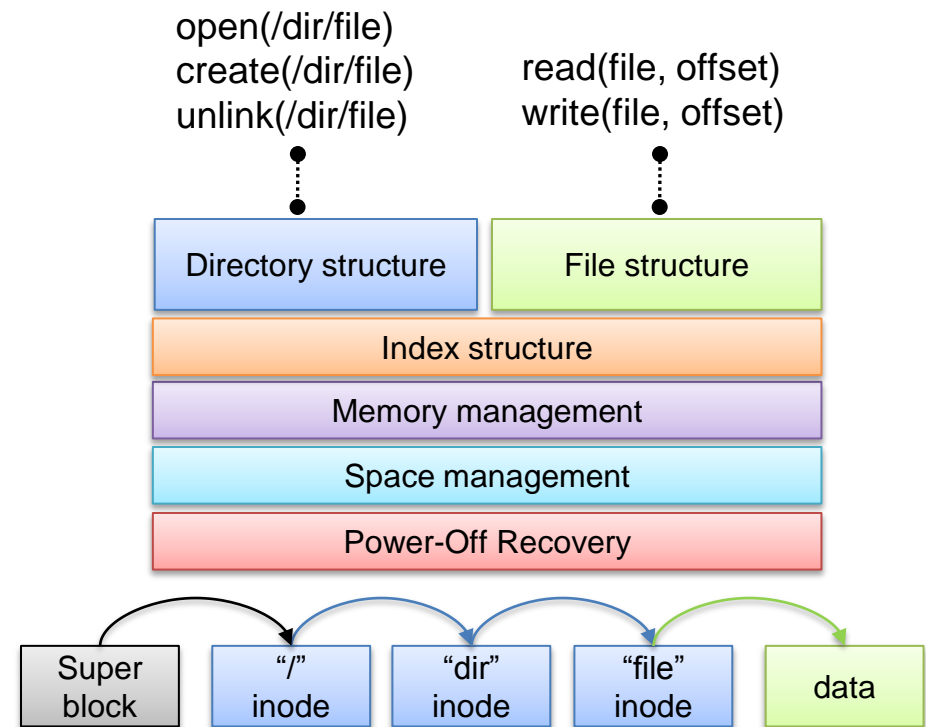

- NAND Flash Memory
  - Erase-before-write
  - Sequential writes inside the erase unit
  - Limited program/erase (P/E) cycle
- Flash Translation Layer (FTL)
  - Garbage collection
  - Wear-leveling
  - Bad block management
- Host-side Issues
  - Poor random write performance
  - Life span and reliability



Source: INDILINX Barefoot controller

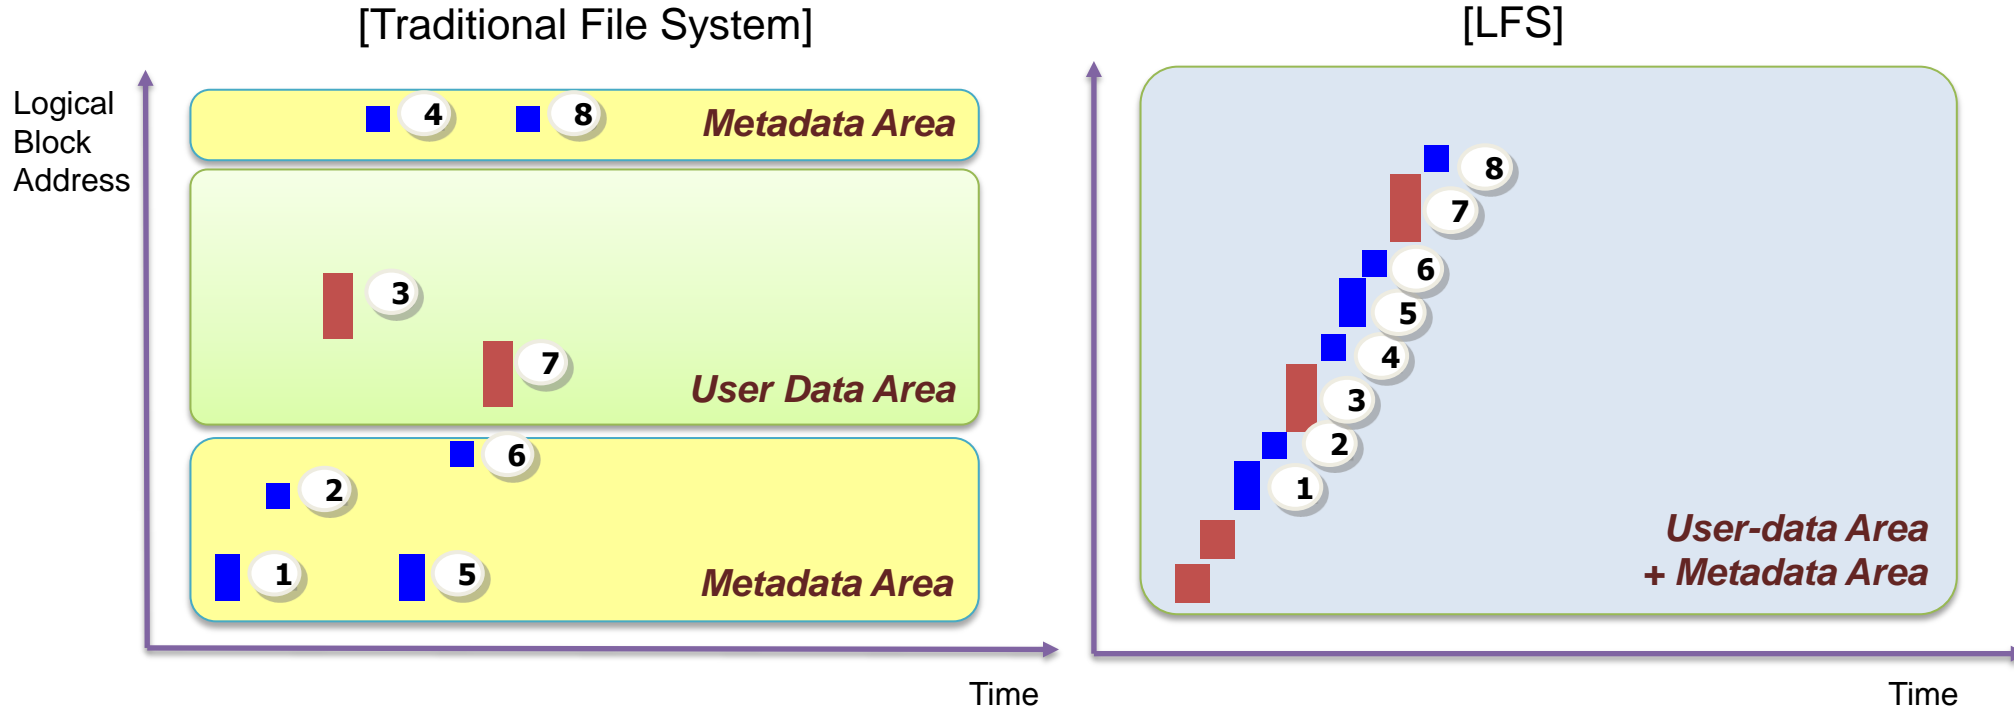
# Introduction

- File System
  - Serve directory and file operations to users
  - Manage the whole storage space



# Log-structured File System

- Log-structured File System (LFS)<sup>[1]</sup>
  - Assume the whole disk space as a big contiguous area
  - Write all the data sequentially
  - Recover quickly with “checkpoint”



[1] Mendel Rosenblum and John K. Ousterhout. 1992. The design and implementation of a log-structured file system. *ACM Trans. Comput. Syst.* 10, 1 (February 1992), 26-52.

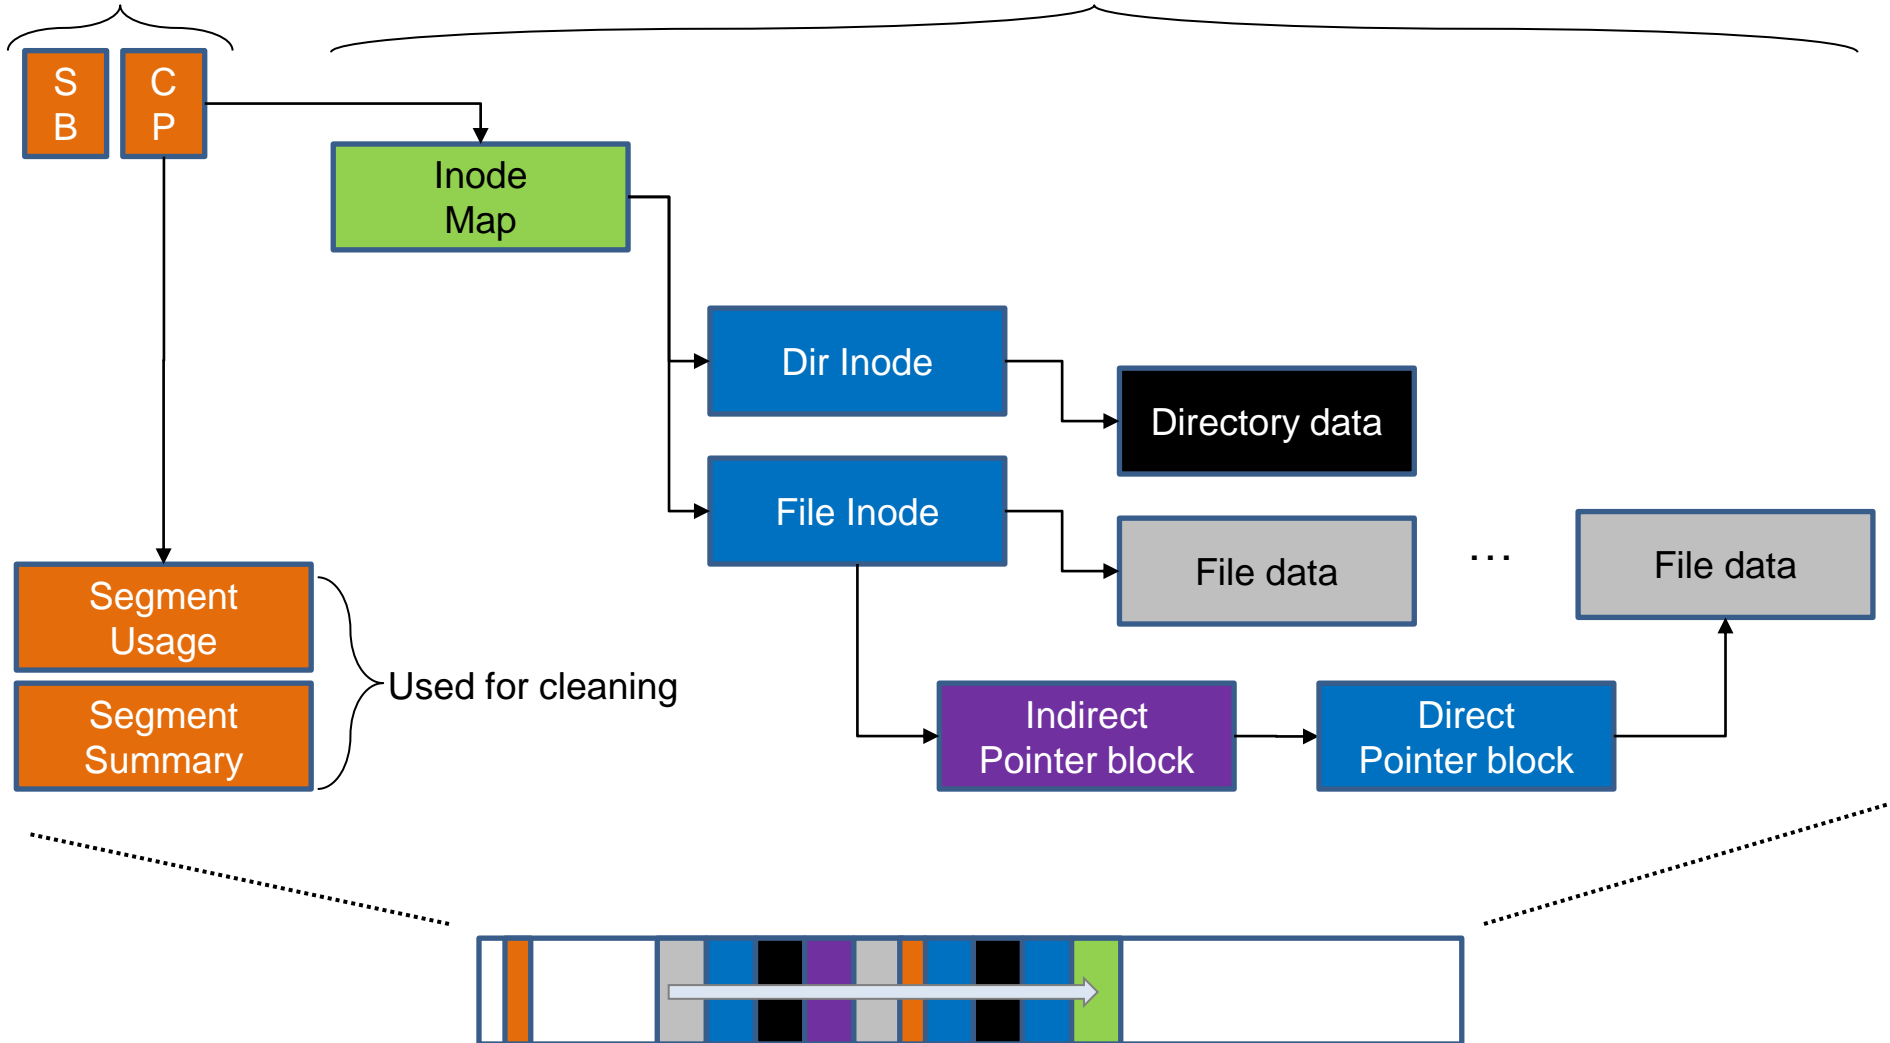


# Log-structured File System (Index Structure)



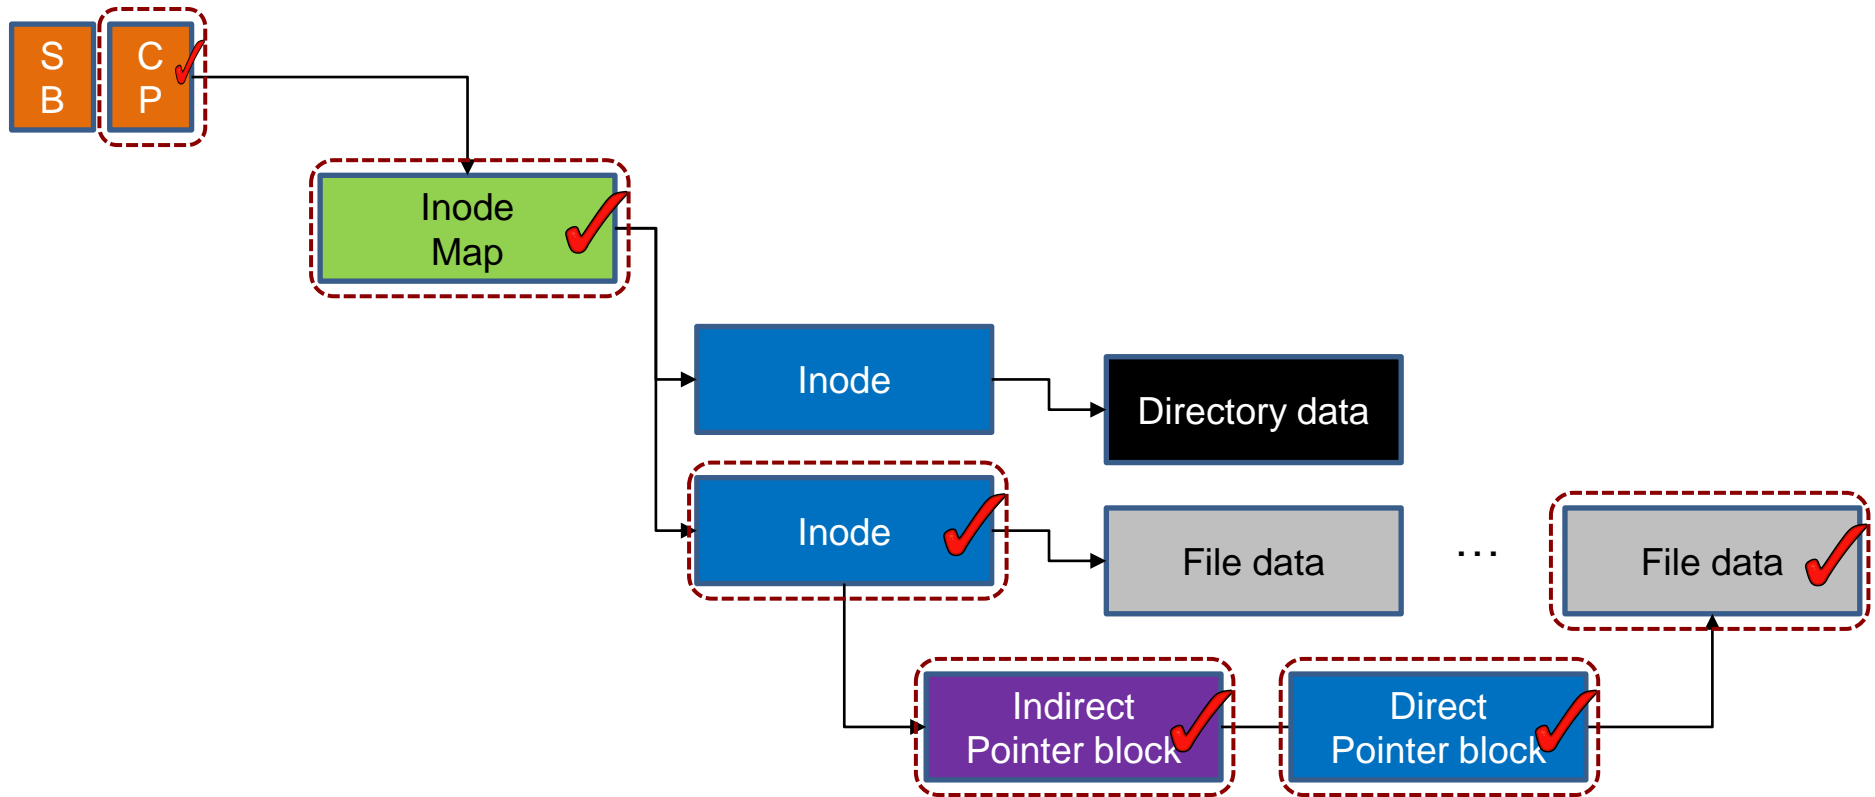
Fixed location, but separated

One big log



# Design Issues I

- Wandering tree problem [2]
  - Propagates index updates recursively
- Goal
  - Eliminate or relax the update propagation



[2] BITYUTSKIY, A. 2005. JFFS3 design issues. <http://www.linux-mtd.infradead.org/>.





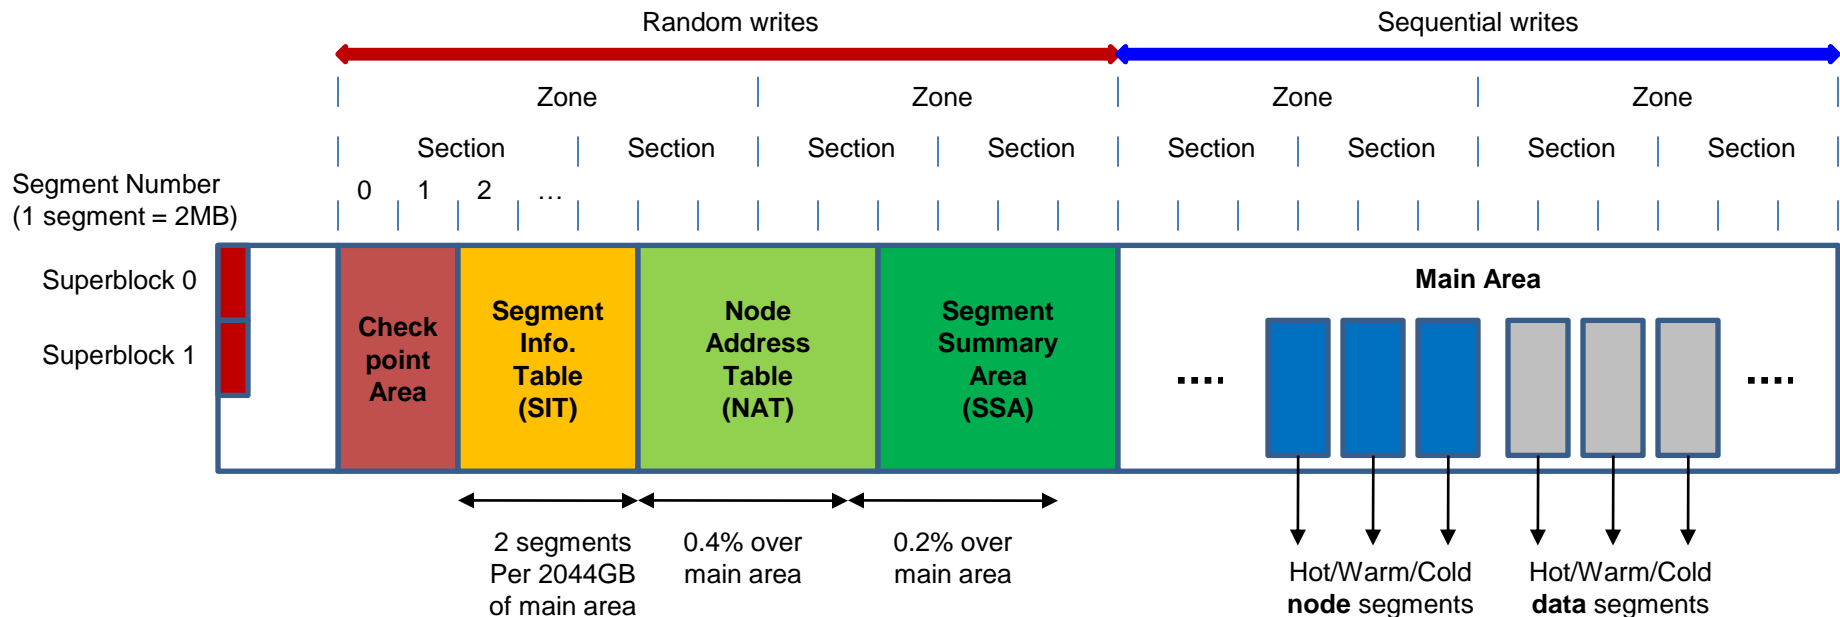
- Cleaning Process
  - Reclaim obsolete data scattered across the whole storage for new empty log space
  - Get victim segments through referencing segment usage table
  - Load parent index structures of there-in data identified from segment summary blocks
  - Move valid data by checking their cross-reference
- Goal
  - Hide cleaning latencies to users
  - Reduce the amount of valid data to be moved
  - Move data quickly
- Specific Issues
  - Cleaning in the background
  - Victim selection policy
  - Hot and cold data separation
  - Instant valid data identification



- Flash Awareness
  - Enlarge the random write area for performance, but provide the high spatial locality
  - Align FS data structures to the operational units in FTL
- Wandering tree problem
  - Use a term, “node”, that represents inodes as well as various pointer blocks
  - Introduce Node Address Table (NAT) containing the locations of all the “node” blocks
- Cleaning overhead
  - Support a background cleaning process
  - Support greedy and cost-benefit algorithms for victim selection policies
  - Support multi-head logs for static hot and cold data separation
  - Introduce adaptive logging for efficient block allocation

# Design of F2FS (On-disk Layout)

- Flash Awareness
  - All the FS metadata are located together for locality
  - Start address of main area is aligned to the zone size
  - Cleaning operation is done in a unit of section
- Cleaning overhead
  - Six active logs for static hot and cold data separation

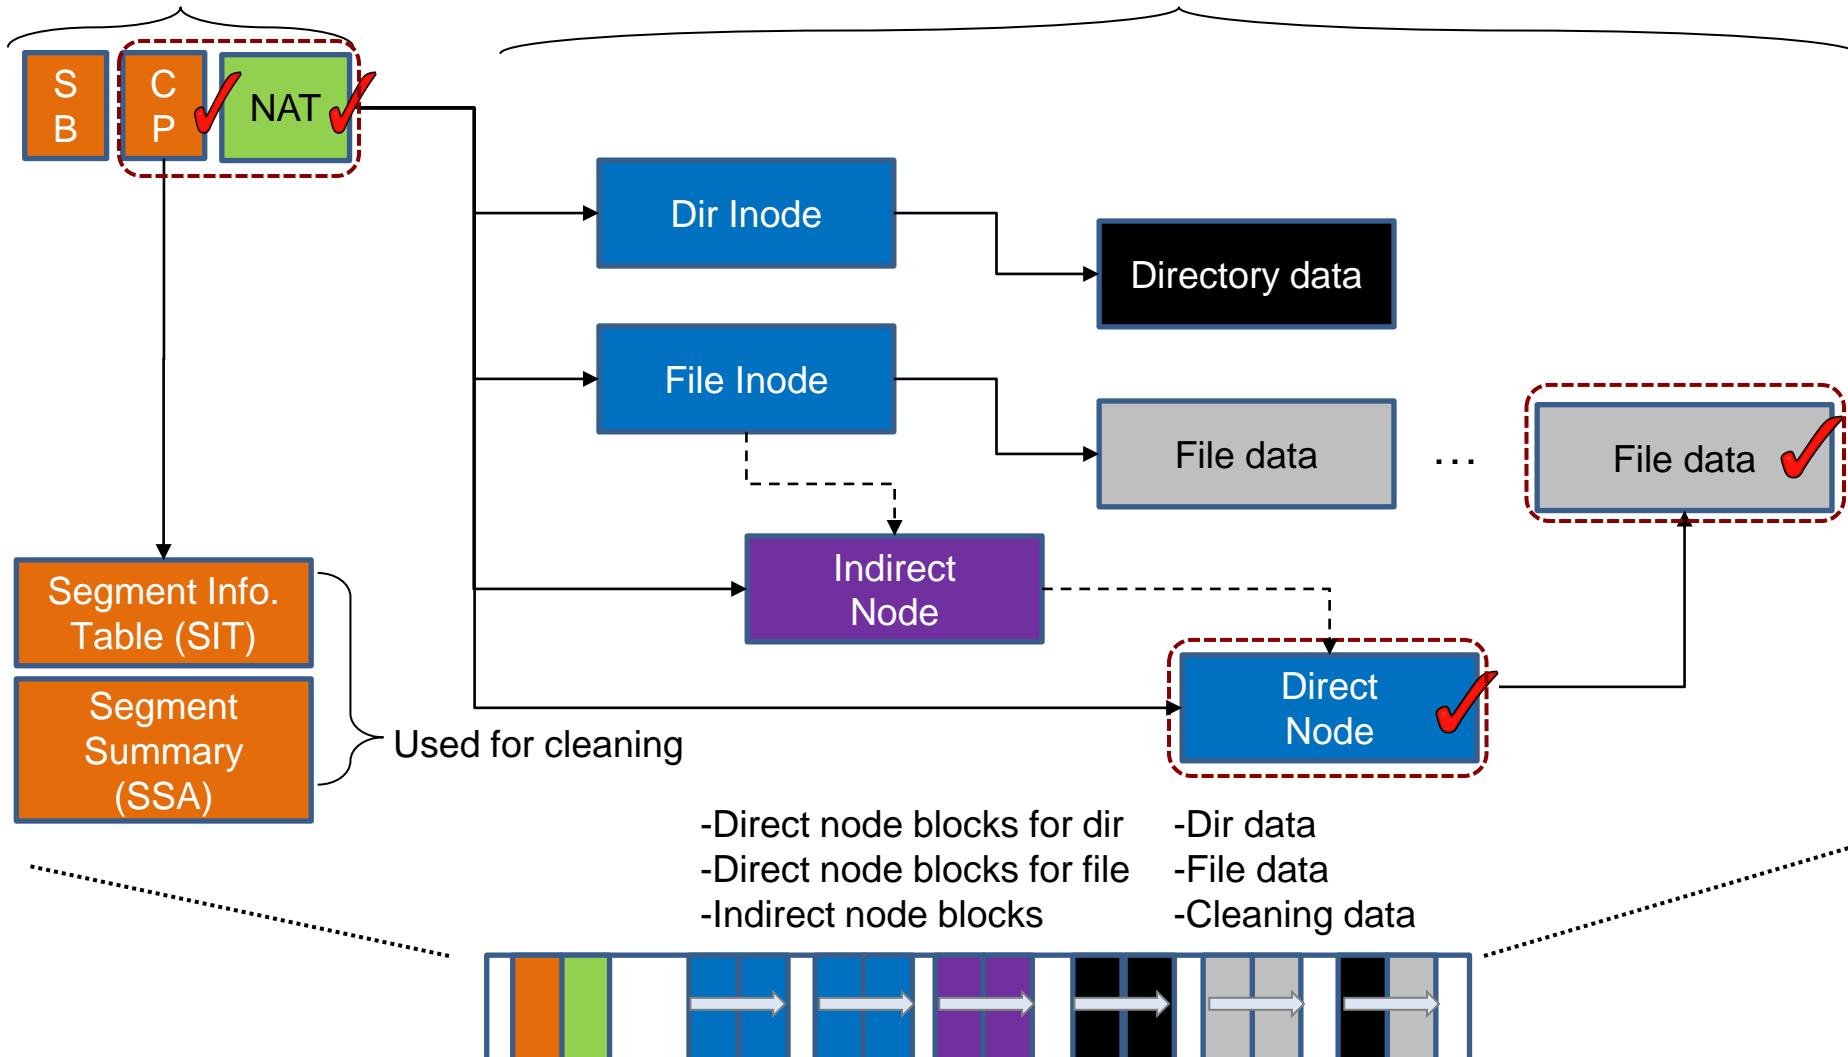


# Design of F2FS (Index Structure)



Fixed location w/ locality

Multiple logs






- Background cleaning process
  - A kernel thread doing the cleaning job periodically at idle time
- Victim selection policies
  - Greedy algorithm for foreground cleaning job
  - Cost-benefit algorithm for background cleaning job
- Block allocation policy
  - Threaded logging
    - Reuse obsolete blocks without cleaning operations
    - Cause random writes
  - Copy-and-compaction
    - Need cleaning operations with some latency
    - Cause no random writes
  - Adaptive logging
    - Normally, copy-and-compaction is adopted
    - If there is not enough free space, the policy is dynamically changed to threaded logging

# Performance Evaluation (micro benchmark)




## [ System Specification ]


CPU	ARM Coretex-A9 1.2GHz
DRAM	1GB
Storage	Samsung eMMC 64GB
Kernel	<b>Linux 3.3</b>
Partition Size	<b>12 GB</b>



## [ iozone ]



## [ fs\_mark ]



## [ bonnie++ ]



# Performance Evaluation (Galaxy S2)



Quadrant	ext4	f2fs	
I/O Performance	<b>3476</b>	<b>3724</b>	<b>248(7.13%)</b>

RLBench	ext4	f2fs	
Overall(sec)	<b>40.76 (0.98)</b>	<b>33.57 (2.01)</b>	<b>-7.18(17.62%)</b>
1000 INSERTs(sec)	<b>20.59 (0.9)</b>	<b>11.96 (1.82)</b>	<b>-8.63(41.93%)</b>
25000 INSERTs in a transaction(sec)	1.79 (0.08)	1.71 (0.01)	-0.08(4.48%)
25000 INSERTs into an indexable table in a transaction(sec)	1.79 (0.08)	1.75 (0.03)	-0.03(1.9%)
100 SELECTs without an index(sec)	0.08 (0.04)	0.05 (0.02)	-0.03(34.21%)
100 SELECTs on a string comparison(sec)	0.07 (0.02)	0.15 (0.21)	0.08(-108.33%)
Creating an index(sec)	0.82 (0.04)	0.94 (0.09)	0.12(-14.08%)
5000 SELECTs with an index(sec)	1.47 (0.11)	1.54 (0.06)	0.07(-4.75%)
1000 UPDATEs without an index(sec)	4.48 (0.04)	4.48 (0.12)	0(0%)
25000 UPDATEs with an index(sec)	3.99 (0.08)	4.14 (0.18)	0.15(-3.81%)
INSERTs from a SELECT(sec)	1.62 (0.15)	1.81 (0.27)	0.19(-11.7%)
DELETE without an index(sec)	1.47 (0.25)	2.02 (0.43)	0.55(-37.41%)
DELETE with an index(sec)	1.43 (0.26)	1.64 (0.3)	0.21(-14.85%)
DROP TABLE(sec)	1.16 (0.11)	1.48 (0.2)	0.31(-26.98%)

Reduce total execution time by 18% over ext4  
Reduce DB insertion time by 42% over ext4

Androbench	ext4	f2fs	
Sequential Read(MB/s)	41.58 (2.72)	41.78 (2.05)	0.2(0.48%)
Sequential Write(MB/s)	4.81 (1.19)	5.63 (1.15)	0.82(17.05%)
Random Read(MB/s)	3.39 (0.06)	3.46 (0.07)	0.07(2.12%)
Random Write(MB/s)	0.25 (0.01)	0.48 (0.01)	<b>0.23(93.5%)</b>
SQLite Insert(s)	15.05 (0.37)	16.63 (0.39)	1.58(-10.5%)
SQLite Update(s)	6.28 (0.27)	3.51 (0.31)	<b>-2.77(44.16%)</b>
SQLite Delete(s)	6.49 (0.19)	3.89 (0.56)	-2.59(39.96%)

Improve random write performance by 94% over ext4

Reduce DB update time by 44% over ext4

# Performance Evaluation (Galaxy S2)

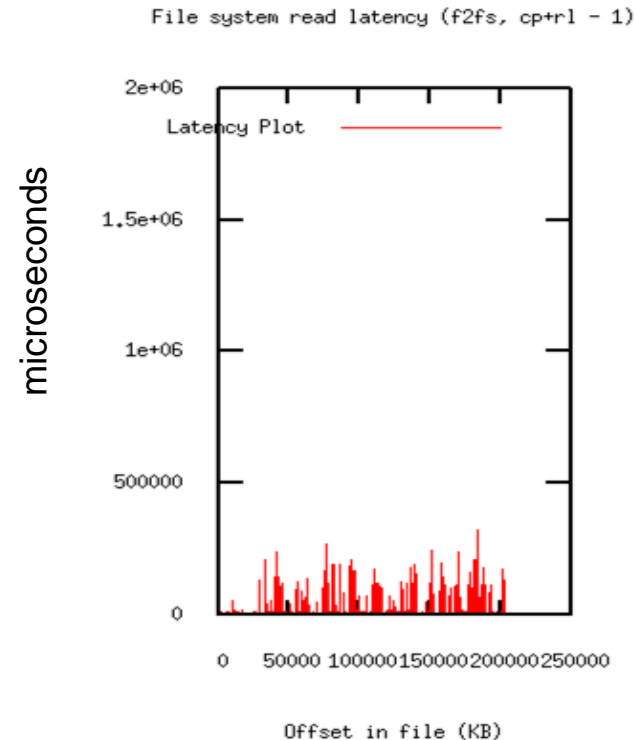



- Multi-tasking Scenario

- Run RLBench and file copying in background
- Measure video streaming latencies in foreground
- High latency may cause cracks in videos

- Results


- Ext4 shows a latency by up to 2.5s, and over 0.5s of latencies more than 10 times
- F2FS shows a latency by up to 0.35s



# Adaptive Logging

- Setup

- Embedded system with eMMC 12GB partition
- I/Ozone random write tests on Several 1GB files
- ext4 shows about 4,000 KB/s continuously




# Victim Selection Policies



- Setup

- In x86, set 3.7 GB partition
- Create three files having 1GB data
- Write 256MB data randomly to the three files
- Write 256MB data randomly to one of them, 30 times



# Galaxy Nexus



CPU	ARM Coretex-A9 1.2GHz
DRAM	1GB
Storage	Samsung eMMC (VFX) 16GB
Kernel	<b>3.0.8</b>
Android ver.	<b>Ice Cream Sandwich</b>

## < Clean >

Items		Ext4	F2FS	Improv.
Contact sync time (seconds)		431	358	20%
App install time (seconds)		459	457	0%
RLBench (seconds)		92.6	78.9	17%
IOZoneWith AppInstall (MB/s)	Write	8.9	9.9	11%
	Read	18.1	18.4	2%

## < Aged >

Items		Ext4	F2FS	Improv.
Contact sync time (seconds)		437	375	17%
App install time (seconds)		362	370	-2%
RLBench (seconds)		99.4	85.1	17%
IOZone With AppInstall (MB/s)	Write	7.3	7.8	7%
	Read	16.2	18.1	12%



- Flash-Friendly File System
  - Focused on Performance and Reliability
  - Not, on new fancy functionalities
- Ubuntu 12.04 LTS
  - Format “/” as F2FS
  - Install & compile kernel & run several applications
- Galaxy S2, S3, and Nexus
  - Format “/data” as F2FS
  - Factory reset & run android apps
- Further Optimization
  - Together!





*Thank you!*